

Ridge regression and LASSO for one variable

In this problem we'll try to understand the effect of regularization on model parameters by considering what happens when $n=p=1$, i.e. we have just one data point and just one predictor. In this case the equations are simple and by plotting the results we can build intuition.

In [1]:

```
1 # Part A
2 import numpy as np
3 import matplotlib.pyplot as plt
4
5 def C_rigid(y, beta, lamda):
6     return (y-beta)**2 + lamda*beta**2;
7
8 betas = np.arange(0.0,3.0, 0.1)
9 plt.plot(betas, [ C_rigid(2,beta,0) for beta in betas ],color="red",linewidth=
10 plt.title("the cost function of Ridge Regression")
11 plt.xlabel("beta")
12 plt.ylabel("C_rigid")
13 plt.show()
```

executed in 2.30s, finished 01:09:20 2018-11-14

<matplotlib.figure.Figure at 0x10b1dfcd0>

What is beta, the value of beta that minimizes $C_{\text{ridge}} = (2-\beta)^2$?

Since the differential of C_{ridge} is $d(C_{\text{ridge}})/d(\beta) = -2(2-\beta)$

so this is minimized when beta is 2



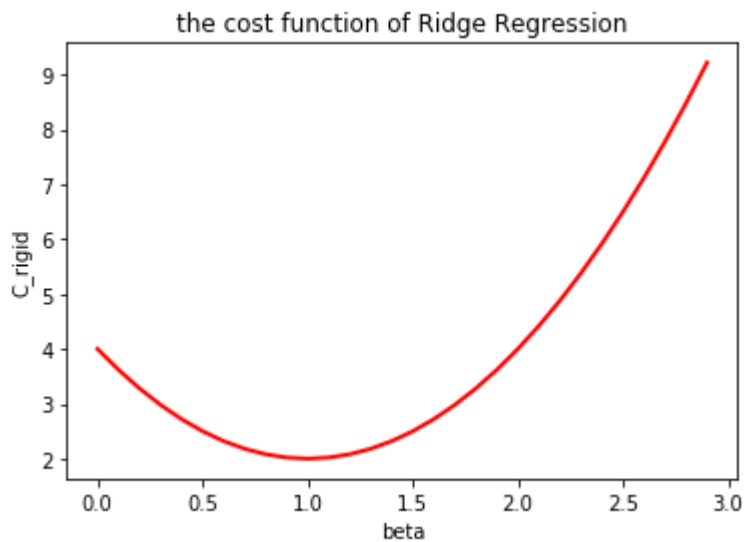
In [2]:

```

1  #part B
2  # Now plot Cridge as a function of  $\beta$  for  $\lambda = 1$ 
3
4  betas = np.arange(0.0,3.0, 0.1)
5  plt.plot(betas, [ C_rigid(2,beta,1) for beta in betas ],color="red",linewidth=
6  plt.title("the cost function of Ridge Regression")
7  plt.xlabel("beta")
8  plt.ylabel("C_rigid")
9  plt.show()

```

executed in 387ms, finished 01:09:21 2018-11-14



Show that the new beta is given by the formula in (6.14), . Describe, in words, how changes as $\beta^* / (1) / 2$ $\beta^* = y + \lambda = y \beta^* \lambda$ increases -- this illustrates why ridge regression is a form of shrinkage .

The new beta that min the cost function is beta=1

since now $C_{ridge} = (2 - \beta)^2 + \gamma \cdot \beta^2$

$d(C_{ridge})/d(\beta) = -2(2 - \beta) + 2 \cdot \beta$

Now if $\gamma = 1$ then $\beta = y/2$

As γ increase the beta will decrease and shrink to 0

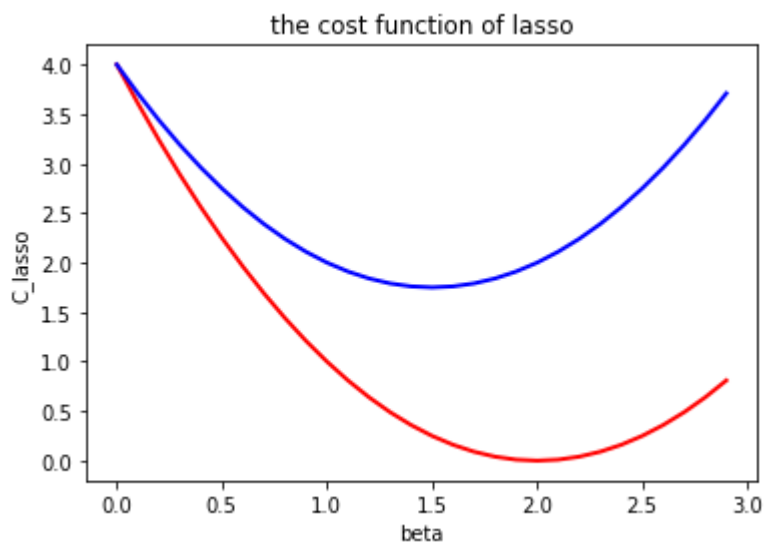
In [3]:

```

1  # part C
2  # Consider the cost function for LASSO, with p=1:
3  # Classo = (y - β)² + λ |β|
4  # Make plots showing Classo as a function of β for λ = 0 and λ = 1 .
5
6
7  def C_lasso(y, beta, lamda):
8      return (y-beta)**2 + lamda*abs(beta);
9
10 betas = np.arange(0.0,3.0, 0.1)
11 plt.plot(betas, [ C_lasso(2,beta,0) for beta in betas ],color="red",linewidth=
12 plt.plot(betas, [ C_lasso(2,beta,1) for beta in betas ],color="blue",linewidth=
13 plt.title("the cost function of lasso")
14 plt.xlabel("beta")
15 plt.ylabel("C_lasso")
16 plt.show()
17

```

executed in 328ms, finished 01:09:21 2018-11-14



The min value of beta for lamda=0 is beta=2 when Lamda=0 it shows the cost func without regularization

The min value of beta for lamda=1 is beta=1.5

Notice that the two cost functions converge at beta=0, which is a result of weight shrinkage.

Forward stepwise model selection

In [4]:

```
1 import matplotlib.pyplot as plt
2 import pandas as pd
3 import numpy as np
4
5 #read data
6 df = pd.read_csv('./anesthesia.csv')
7 data = df.values
8 print data.shape
```

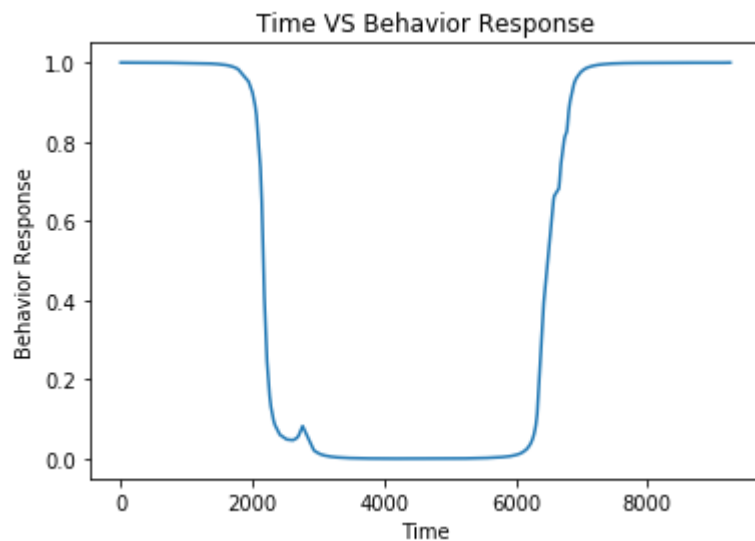
executed in 3.27s, finished 01:09:24 2018-11-14

(926, 105)

In [20]:

```
1 # Part A
2 # Make a plot of Time Vs. BehaviorResponse for the anesthesia data.
3 time = data[:,0]
4 BehaviorResponse = data[:,104]
5 plt.plot(time, BehaviorResponse)
6 plt.title("Time VS Behavior Response")
7 plt.xlabel("Time")
8 plt.ylabel("Behavior Response")
9 plt.show()
10
```

executed in 250ms, finished 01:23:39 2018-11-14




In [35]:

```

1 # Part b. Make a scatter plot of BehaviorResponse vs. EEG power at 0Hz
2 plt.plot(df['F0Hz_1'], df['BehaviorResponse'], 'ro');
3 plt.ylabel('Behavior Response');
4 plt.xlabel('Power at 0Hz');

```

executed in 287ms, finished 01:34:19 2018-11-14



Describe, in words, the relationship between these variables.

As EEG power increases at 0H, the prob of BehaviorResponse=1 decreases

In [41]:

```

1 ## Part C
2 #What is the correlation coefficient between Behavior Response and EEG power a
3 import numpy as np
4
5 corr_coef = np.corrcoef(df['F0Hz_1'], df['BehaviorResponse']);
6 print 'The correlation coefficient is ', corr_coef[0][1]

```

executed in 9ms, finished 01:37:17 2018-11-14

The correlation coefficient is -0.4632119330934692

In [52]:

```

1  # Part D
2  # Fit a simple linear regression model of the form, BehaviorResponse ~ 1 + F0Hz_1
3  import numpy as np
4  import pandas as pd
5  import matplotlib.pyplot as plt
6  import statsmodels.api as sm
7  from patsy import dmatrices
8
9  ### do linear regression
10 y, X = dmatrices('BehaviorResponse ~ 1 + F0Hz_1', data=df, return_type='dataframe')
11 X = sm.add_constant(X)
12 mod = sm.OLS(y, X)
13 # fit model
14 res = mod.fit()
15 # look at results
16 print(res.summary())

```

executed in 54ms, finished 01:55:28 2018-11-14

OLS Regression Results

```

=====
=====
Dep. Variable:          BehaviorResponse    R-squared:
    0.215
Model:                  OLS    Adj. R-squared:
    0.214
Method:                 Least Squares    F-statistic:
    252.4
Date:                  Wed, 14 Nov 2018    Prob (F-statistic):
1.96e-50
Time:                  01:55:28    Log-Likelihood:
-510.49
No. Observations:      926    AIC:
    1025.
Df Residuals:          924    BIC:
    1035.
Df Model:               1

Covariance Type:      nonrobust

=====
=====

```

	coef	std err	t	P> t	[0.025
Intercept	1.8503	0.084	22.000	0.000	1.685
F0Hz_1	-0.4433	0.028	-15.888	0.000	-0.498

```

-----
-----
Omnibus:              995.624    Durbin-Watson:
    0.198
Prob(Omnibus):        0.000    Jarque-Bera (JB):
    63.317
Skew:                 -0.152    Prob(JB):
1.78e-14
Kurtosis:             1.755    Cond. No.

```

20.3

```
=====
```

Warnings:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

Is the slope parameter statistically significant?

Since the p-value is smaller than 0.05, there is a statistically significant relationship between BehaviorResponse and power at 0Hz

In [122]:

```
1  ▾ # part E
2    # Fit a multiple linear regression that uses all of the EEG power features
3
4    features = df.head(0)
5    features_str = '1 '
6  ▾ for feature in features:
7  ▾     if feature != "Time" and feature != "BehaviorResponse":
8         features_str += ' + ' + feature;
```

executed in 7ms, finished 03:23:01 2018-11-14

In [127]:

```

X1= df[['F0Hz_1', 'F125Hz_103']]
y2= df['BehaviorResponse']
3
### do linear regression
y5 X = dmatrices('BehaviorResponse ~' + features_str, data=df, return_type='dataframe')
X6= sm.add_constant(X)
mod = sm.OLS(y, X)
#fit model
res = mod.fit()
#look at results
print(res.summary())

```

executed in 697ms, finished 03:24:10 2018-11-14

Intercept	0.6519	0.131	4.994	0.000	0.396
0.908					
F0Hz_1	-0.0122	0.017	-0.698	0.485	-0.046
0.022					
F1Hz_2	-0.0394	0.025	-1.560	0.119	-0.089
0.010					
F3Hz_3	0.0006	0.028	0.020	0.984	-0.054
0.055					
F4Hz_4	-0.0060	0.027	-0.226	0.822	-0.058
0.046					
F5Hz_5	0.0008	0.025	0.032	0.974	-0.049
0.050					
F6Hz_6	0.1510	0.022	6.858	0.000	0.108
0.194					
F8Hz_7	-0.0003	0.024	-0.012	0.991	-0.047
0.047					
F9Hz_8	0.0066	0.027	0.247	0.805	-0.046
0.059					
F10Hz_9	-0.1492	0.024	-6.324	0.000	-0.195
-0.103					

What is the p-value of the slope for F0Hz_1?

0.485

Is it statistically significant?

The slope coefficient for power at F0Hz_1 is Not significant Now

In [206]:

```
1  ▼ # part F
2  # Fit a 103 models and find the single features with the lowest MSE
3  # for each fit use the MSE
4  # make a plot showing MSE VS feature number
5
6  mse_list = [0 for i in range(103)]
7  feature_list = []
8  ▼ for feature in features:
9  ▼     if feature != "Time" and feature != "BehaviorResponse":
10         feature_list.append(feature)
11
12     # 103 iterations
13     index = 0
14  ▼ for feature in feature_list:
15         y, X = dmatrices('BehaviorResponse ~' + feature, data=df, return_type='dat
16         X = sm.add_constant(X)
17         mod = sm.OLS(y, X)
18         # fit model
19         res = mod.fit()
20         #predict the data
21         y_hat = res.predict(X);
22         mse_list[index] = np.mean( (df['BehaviorResponse'] - y_hat)**2 )
23         index += 1
24
```

executed in 1.19s, finished 04:19:36 2018-11-14

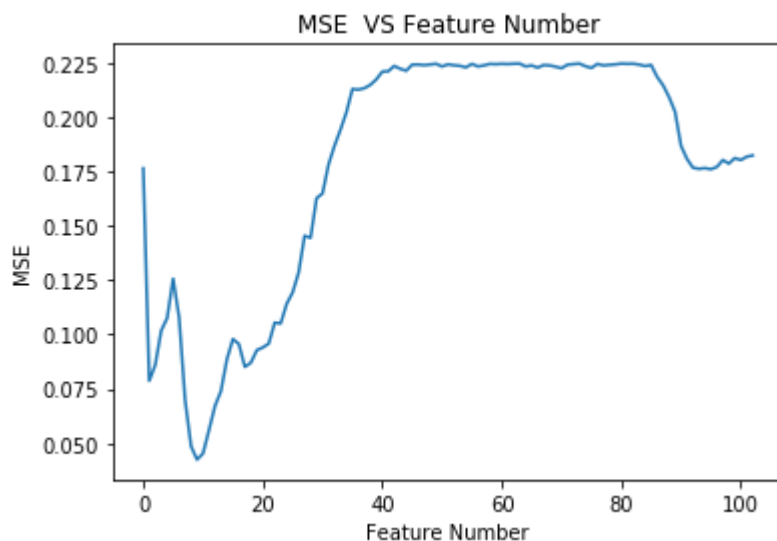
In [207]:

```
1 print min(mse_list)
2 print "So F11Hz_10 is winning feature"
3
4 x=range(103)
5 y= []
6 for index in range(103):
7     y.append( mse_list[index] )
8
9 plt.plot(x,y)
10 plt.title("MSE VS Feature Number")
11 plt.xlabel("Feature Number")
12 plt.ylabel("MSE")
13 plt.show()
```

executed in 281ms, finished 04:19:37 2018-11-14

0.0423969081374

So F11Hz_10 is winning feature



In [213]:

```

1  # Part G
2  # Now write a loop to fit 102 models of the form, BehaviorResponse ~ 1 + X1 +
3  # the best feature obtained from part (f) and X2 is one of the other features.
4  # Which combination of two features gives the best prediction?
5
6  best_feature = "F11Hz_10"
7
8  #build a new list
9  new_feature_list = []
10 for feature in features:
11     if feature != "Time" and feature != "BehaviorResponse" and feature!="F11Hz
12         new_feature_list.append(feature)
13
14 print len(new_feature_list)
15
16 # 102 iterations
17 index = 0
18 mse_list = [0 for i in range(103)]
19 for feature in new_feature_list:
20     y, X = dmatrices('BehaviorResponse ~ 1 + ' + best_feature \
21                     + '+' + feature, data=df, return_type='dataframe')
22     X = sm.add_constant(X)
23     mod = sm.OLS(y, X)
24     # fit model
25     res = mod.fit()
26     #predict the data
27     y_hat = res.predict(X);
28     mse_list[index] = np.mean( (df['BehaviorResponse'] - y_hat)**2 )
29     index += 1
30
31 print 'the winning combination is ', best_feature, \
32       ' and ', new_feature_list[9]

```

executed in 1.25s, finished 04:24:05 2018-11-14

102

the winning combination is F11Hz_10 and F12Hz_11

In [241]:

```

1  ▾ # Part H
2  from sklearn.model_selection import KFold
3
4  best_comb = "F11Hz_10 + F12Hz_11";
5  Num_Splits = 10
6  Num_observation = len(df['BehaviorResponse'])
7  kf = KFold(n_splits=10)
8  X_Splits = kf.get_n_splits(Num_observation)
9  MSE_matrix = ones(10,2)
10
11  #iterate 10 times
12  ▾ for train_index, test_index in kf.split(X):
13      train = (kf == index)
14      test = test_ind
15      y, X = dmatrices('BehaviorResponse ~ 1 + ' + best_comb, return_type='dataf
16      X = sm.add_constant(X)
17      mod = sm.OLS(y, X)
18      res = mod.fit()
19      y_hat = res.predict(X);
20      mse_list[index] = np.mean( (df['BehaviorResponse'] - y_hat)**2 )
21
22      model = fitlm(data(train_ind,:), ['BehaviorResponse ~ 1 + ' winning_comb])
23      y_hat = predict(model, data);
24
25      MSE_matrix[index][0] = np.mean( (df['BehaviorResponse'] - y_hat)**2 )
26      MSE_matrix[index][1] = np.mean( (df['BehaviorResponse'] - y_hat)**2 )
27
28  print "So training and testing MSE for this model are printed below"
29  print "train    Test"
30  ▾ for item in MSE_matrix:
31      print item

```

executed in 30ms, finished 04:57:36 2018-11-14

So training and testing MSE for this model are printed below

```

train    Test
[0.037099 0.024939]
[0.035453 0.039801]
[0.036064 0.034278]
[0.036726 0.028319]
[0.036026 0.034621]
[0.034259 0.050515]
[0.036244 0.032711]
[0.036189 0.033118]
[0.035931 0.035774]
[0.034742 0.046127]

```

In []:

1